

Advanced Models II

Practical Machine Learning (with R)

UC Berkeley

Fall 2015

Topics

- Review and Expectations
- Questions
- New Topics



Remaining Topics

Today

- Custom loss functions
- Boosting
- Deployment
- Neural Networks

Next Time:

- svm
- regularization/shrinkage : ridge and lasso
- strategies and “patterns”





REVIEW AND EXPECTATIONS



REVIEW

⇒ Trees

Tree Variants

Surrogate Splits (missing Data)

⇒ Rules

(Related to Trees, relaxes the Path constraint)

⇒ Treatment of Categorical Variables

- Grouped vs Independent

⇒ Tuning Parameters



TWO BIG IDEAS

➤ **Wisdom of the crowds**

It is better to make estimates from multiple models (**ensembles**) than individual models

- Better predictions
- Lower variance for the same model

➤ **Greed is bad. Patience is good.**

It is better to slowly approach your solution than arrive at an answer directly



MODEL IMPROVEMENT

- ➔ Ensembles / Model Averaging
Effect: reduce variance of model
 - M5 Trees (Regression in nodes)
 - Bagging (Bootstrap Aggregation, any learner)
 - Random Forest
 - Advantage / Disadvantages



RANDOM FOREST ADVANTAGES

- ➔ No overfitting
- ➔ More trees better (limited by computation time/power only)
- ➔ In caret, parameters are considered independently
- ➔ Because each learner is selected independently of all previous learners, randomforests is robust to a noisy response
- ➔ Computationally efficient -- each tree built on subset of predictors at each split.
- ➔ Use any tree variants as "base learner": CART, ctree, etc



CUSTOM LOSS/COST FUNCTIONS

- ➔ Most methods use “0-1” or “simple” loss, where all errors treated equal
 - Overly simple
- ➔ Real world

	0	L_{FN}	Binary/two-class
L_{FP}	0	0	
0	ϵ_{12}	ϵ_{13}	N-class
ϵ_{21}	0	ϵ_{23}	
ϵ_{31}	ϵ_{32}	0	

- ➔ Mostly applies to classification problems
- ➔ Not to be confused with error metric for training



NATIVE COST/LOSS MATRIX

Used during training

- For example

```
rpart( parms = list(loss=L) )
```

- ➔ **Best results:** all aspects optimization are optimized through the custom costs
- ➔ Unfortunately, not every methods support custom loss functions



CUSTOM LOSS:

CLASS PROBABILITIES

- ⇒ Use class probabilities ... adjust class based on probabilities
- ⇒ Difficult/time consuming to implement in practice




TWO-CLASS CLASSIFICATION PROBLEM

- ⇒ Use case weights based on the benefit/cost of the correct classification

`Fraud ~ . , weight=dollars`

- Emphasis given to cases that are more important, less likely of erroring here
- Not all methods support case “weights” argument
- Works well for rare-event detection problem
- Lose 0-1 predictability



TWO-CLASS CLASSIFICATION PROBLEM

- ⇒ Define “benefit” against naïve case
 - e.g. assume all transactions are good.

.	F	G
F	0	$-\alpha$
G	$-x$	$+\delta$

outcome	benefit
F	$-x$
T	$+\delta$

- ⇒ Models “expected” benefit from alt. decision
- ⇒ Sign determines classification /action
- ⇒ Does not account for all error terms (e.g. FN)

TWO-STEP MODEL

⇒ Use two models

- **Model 1:** Unmodified Classification model
 - Goal best classification prediction available
 - Assume this is the how the class will be identified
- **Model 2:** Evaluation Model
 - **Model 1** provides both the response and predicted response
 - Calculate benefit (–error) of each poss. class, given the predicted class
 - New Response = class with the highest benefit for each case
- **Benefit:**
Separates classification model from benefit





QUESTIONS



NEW TOPICS



BOOSTING

⇒ Single models work;

- Multiple models work better

⇒ Idea is simple:

- **Fit first** model: $\hat{y}_1 \sim f_1(x)$

- **Fit** errors/residuals:
$$\begin{aligned}\hat{y}_2 &= f_2(y - \hat{y}_1) \\ &= f_2(y - f_1(x)) \\ &= f_2(x)\end{aligned}$$

- **Iterate:** $\hat{y}_i = (y - \hat{y}_{i-1}) \sim f_i(x)$

- **Predict:** $\hat{y} \sim \sum_i f_i(x)$



BOOSTING NOTES

- ➔ Additive models
- ➔ Works best with “weak learners”
 - i.e. ungreedy, low bias, low variance
 - ~~Any~~ Most models with a tuning parameter can be a weak learner
 - Trees are excellent weak learners
 - Weak → “restricted depth”
- ➔ Residuals or errors define a gradient
- ➔ Interpreted as forward step-wise regression with exponential loss



SIMPLE GRADIENT BOOSTING

- 1 Select tree depth, D , and number of iterations, K
- 2 Compute the average response, \bar{y} , and use this as the initial predicted value for each sample
- 3 for $k = 1$ to K do
 - 4 Compute the residual, the difference between the observed value and the *current* predicted value, for each sample
 - 5 Fit a regression tree of depth, D , using the residuals as the response
 - 6 Predict each sample using the regression tree fit in the previous step
 - 7 Update the predicted value of each sample by adding the previous iteration's predicted value to the predicted value generated in the previous step
- 8 end

Simple Gradient Boosting – Comparison To Random Forest

Similarities

Differences



STOCHASTIC GRADIENT BOOSTING

⇒ Gradient Boosting Susceptible to Overfitting

- Apply “regularization/shrinkage”
 - Use λ (“Learning Rate”)
Rather than add the entirety of the residuals, add a fraction of the residuals at each iteration.

$$\hat{y} \sim \lambda \sum_i f_i(x) \quad 0 < \lambda \leq 1$$

- Small values for λ (~ 0.01) work best
- $\lambda \sim 1/\text{computational time} \sim 1/\text{storage time}$

⇒ Use bagging, as well

- Bagging Fraction: a sample of data in each loop iteration

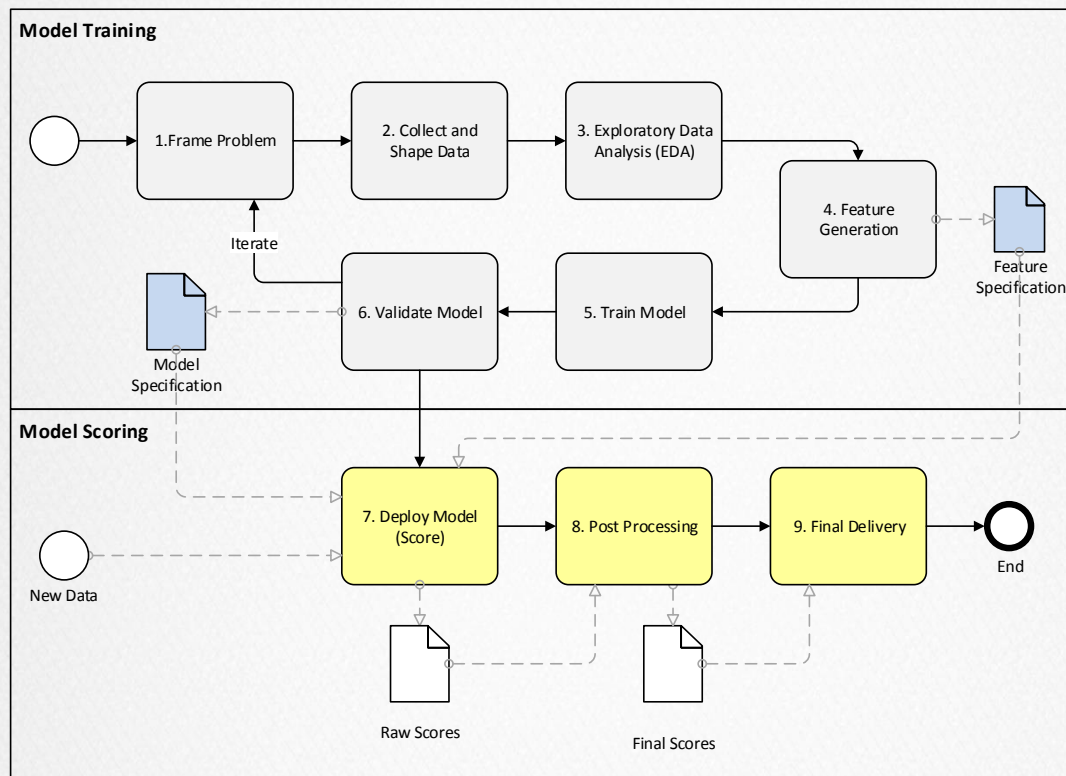


DEPLOYMENT



Deployment

- Making a model available to scoring to a larger group of users



General Consideration

Users	How many? Technical proficiency
Data	Single case < set of cases < universe
Frequency	How often scores accessed?
Latency	How much time tolerable to score?
Interface	Web UI (shiny) Command-line(optigrab) Rserve OpenCPU
Resources Req'd	Memory, disk

General Consideration

- ⇒ Three steps
 - accept inputs from interface
 - apply logic
 - present / render findings
- ⇒ Accept inputs /render findings
 - Dependent upon Interface
 - R / command line / shiny / web application



General Considerations : Application Logic

- ⇒ Training / scoring data must be ***similar***
 - Every transformation made to create training data must be “replayed” on newdata
 - scale must use center and scale parameters from training
 - impute must use same distributions/data
 - Etc.
- ⇒ standardize feature development
 - [fetch data] → build features → build frame
 - Separate functions



Native R

- ➔ Most flexible
- ➔ Requires R knowledge, tech. proficiency

➔ Best Practices

- standardize model location and usage
- use R package features(?)

- `data`
- `fetch`
- `featurize`
- `build_frame`

Deploying of model is simple as `library(mypackage)`



Command-line application

users access models from command prompt

```
> Rscript score.R --sepal.length 2.7 -sepal.width
```

Assumptions

- No specific R knowledge required
- Some technical proficiency
- Data passed in as **optigrab**
 - Actions bound to parameters
- Best with single-case scoring
- Can work with data sets if you can specify them



Shiny

⇒ *Key Features*

- **Reactive** programming

- Variable have dependency on other values
- Update values when dependencies change

- Separated concerns

- **ui/ui.R** (presentation) : high-level functions for widgets and layout
- **server / server.R** (application) : application logic



ASSIGNMENT



EVALUATING VARIABLE CONTRIBUTIONS IN COMPLEX MODELS

- How do understand the contribution of each variable?
- Retrain the model leaving out the variable: change is the contribution
- Randomly permit values in evaluation set and compare the performance (Brieman)
- Aggregate performance of variable across each model

APPENDIX



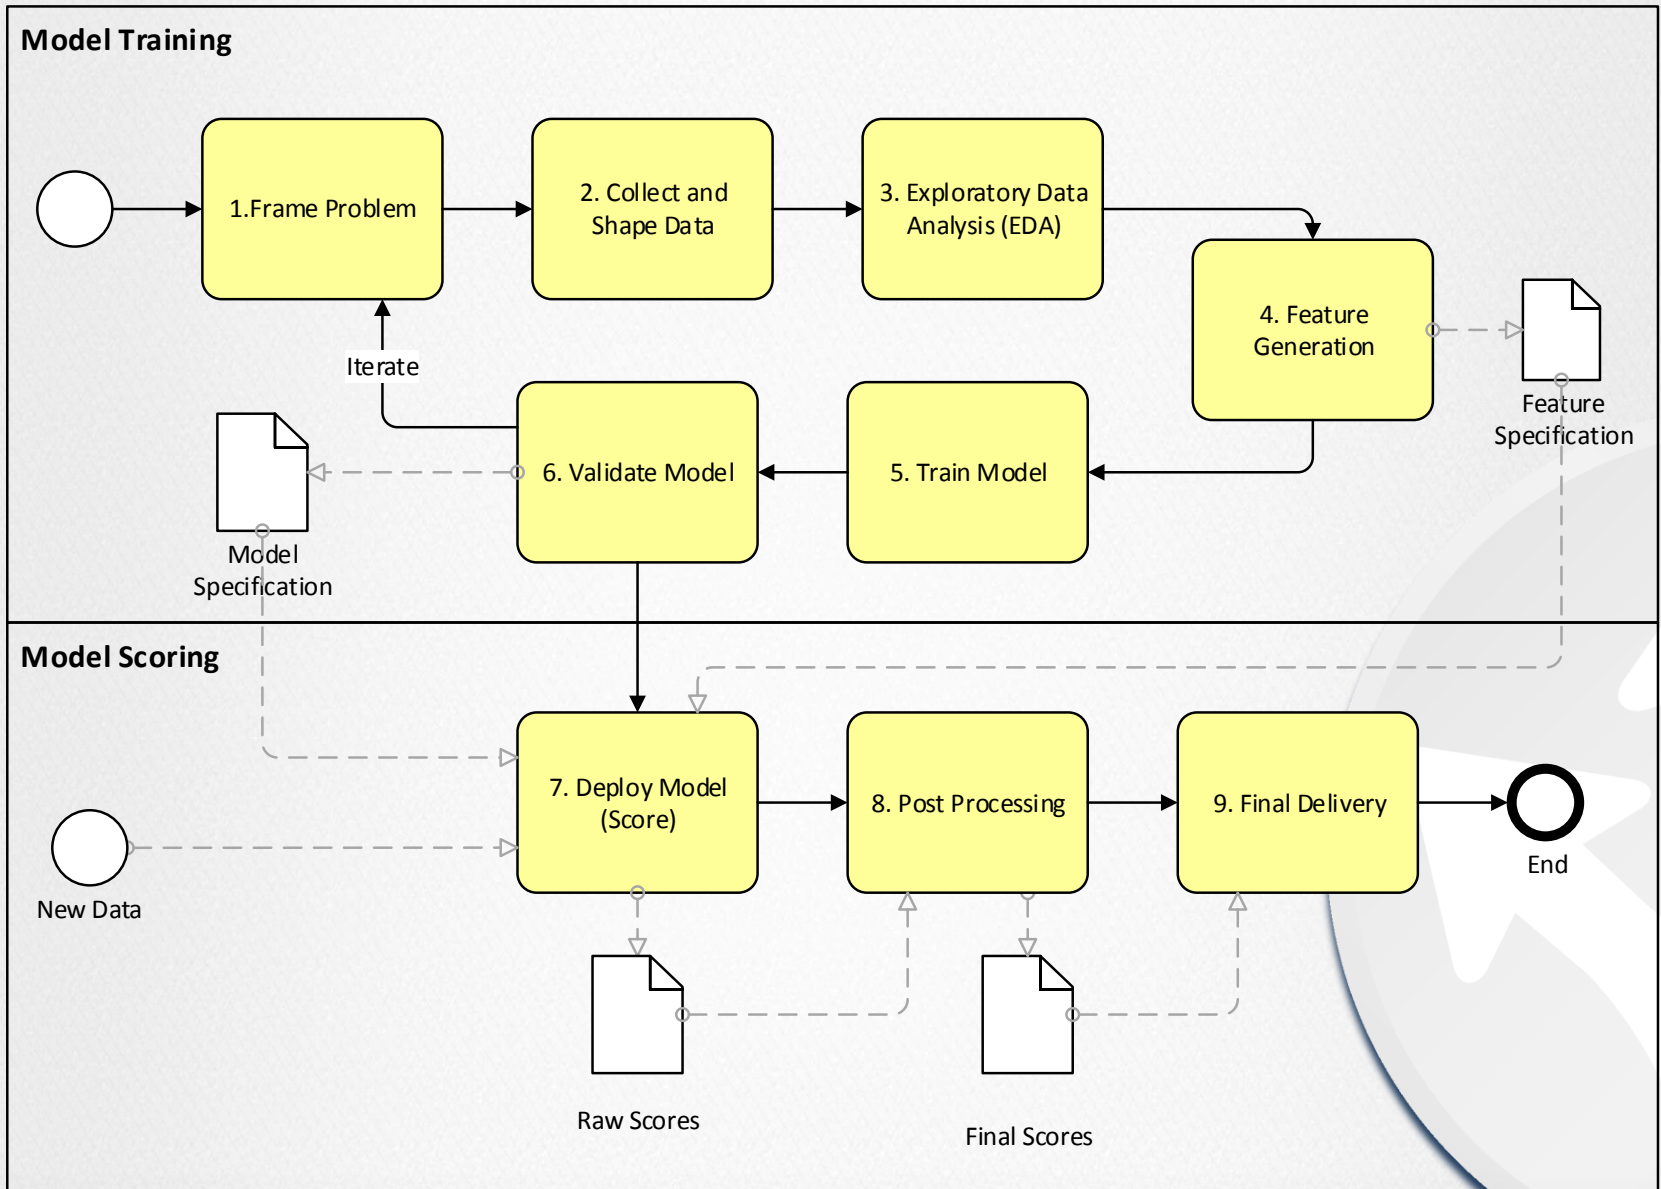
ADABOOST (SHAPIRE/FREUND)

- 1 Let one class be represented with a value of $+1$ and the other with a value of -1
- 2 Let each sample have the same starting weight ($1/n$)
- 3 **for** $k = 1$ **to** K **do**
- 4 Fit a weak classifier using the weighted samples and compute the k th model's misclassification error (err_k)
- 5 Compute the k th stage value as $\ln((1 - err_k) / err_k)$.
- 6 Update the sample weights giving more weight to incorrectly predicted samples and less weight to correctly predicted samples
- 7 **end**
- 8 Compute the boosted classifier's prediction for each sample by multiplying the k th stage value by the k th model prediction and adding these quantities across k . If this sum is positive, then classify the sample in the $+1$ class, otherwise the -1 class.

MULTI-CLASS PERFORMANCE



Comprehensive ML Process



TERMS

- ⇒ Kappa Statistic,
 - ⇒ S-Statistics, F-Statistic
-



EXAMPLE OF ML ALGORITHM(S)

- ➔ Spam Filter
- ➔ handwriting recognition (svm)
- ➔ Traffic engineering (lights)
- ➔ Weather prediction
- ➔ Sentiment analysis (social media)
- ➔ Netflix Recommender
- ➔ Fraud detection (Visa)
- ➔ Imaging processing
- ➔ (network) Intrusion detection
- ➔ Self-driving cars



COMPARISON OF MODELS (CHART)

